

Question 4

You have recently started in the market research division of a box company in Tustin, CA. As one of your first tasks, Randy, the head of market research, asks you to analyze the following data from a recent poll of 5,000 Californians:

Question: If you were to purchase boxes in the next few months, would you prefer corrugated or uncorrugated cardboard for your boxes?

| | | | California Population | Respondents | Prefer Corrugated |
|--------|------------|------------|--------------------------|-------------|----------------------|
| gender | age_bucket | urbanicity | | | |
| male | 18_to_35 | urban | 4,815,108 | 252 | 164 |
| female | 18_to_35 | urban | 4,151,623 | 260 | 145 |
| male | 18_to_35 | rural | 2,342,416 | 234 | 156 |
| female | 18_to_35 | rural | 1,854,720 | 228 | 123 |
| male | 36_to_64 | urban | 6,676,992 | 678 | 451 |
| female | 36_to_64 | urban | 6,259,680 | 714 | 373 |
| male | 36_to_64 | rural | 3,338,496 | 684 | 434 |
| female | 36_to_64 | rural | 2,543,616 | 568 | 293 |
| male | 65_plus | urban | 1,516,234 | 354 | 225 |
| female | 65_plus | urban | 1,608,023 | 428 | 224 |
| male | 65_plus | rural | 741,888 | 260 | 168 |
| female | 65_plus | rural | 927,360 | 341 | 187 |

For your convenience, the individual-level responses from the poll are provided in "poll_responses.csv" which you should have received along with this exam (yes responses are coded as 1 and no responses as 0 in the file). The questions below can be answered with or without this additional data.

Part A. What is your best estimate for the percentage of Californians who prefer corrugated cardboard for their boxes? What is the 95% confidence interval for this estimate? Please show your work.

```
In [13]: import statsmodels #you seem to need this AND the next line
#from statsmodels.stats.proportion import proportions_ztest
import pandas as pd
```

```
In [8]: poll = pd.read_csv("poll_responses.csv")
poll.head(2)
```

Out[8]:

| | gender | age_bucket | urbanicity | prefer_corrugated |
|---|--------|------------|------------|-------------------|
| 0 | male | 18_to_35 | urban | 1 |
| 1 | male | 18_to_35 | urban | 1 |

```
In [9]: x = poll.prefer_corrugated
x.head(4) #need a few more rows to properly see it
```

Out[9]:

| | |
|---|---|
| 0 | 1 |
| 1 | 1 |
| 2 | 1 |
| 3 | 1 |

Name: prefer_corrugated, dtype: int64

```
In [14]: count = x.sum()
nobs = x.count()
statsmodels.stats.proportion.proportion_confint(\
    count, nobs, alpha=0.05, method='normal')
#alpha is significance level; #method is normal for Z test
#lower bound, upper bound of 95% confidence interval
```

Out[14]:

| | |
|----------------------|---------------------|
| (0.5748433633316954, | 0.6021212437468888) |
|----------------------|---------------------|

Part B. Are men and women significantly different in their likelihood to prefer corrugated cardboard?

First, we will solve this with a z test, since that's the way I did it on the exam. Then we will solve it with a t test, since that can also be permissible.

Two sample proportion z test

```
In [11]: #first we start to do the analogy to a SQL where clause to get the genders set up
poll[poll["gender"]=="male"].head(3)
```

Out[11]:

| | gender | age_bucket | urbanicity | prefer_corrugated |
|---|--------|------------|------------|-------------------|
| 0 | male | 18_to_35 | urban | 1 |
| 1 | male | 18_to_35 | urban | 1 |
| 2 | male | 18_to_35 | urban | 1 |

```
In [12]: poll[poll["gender"]=="female"].head(3)
```

Out[12]:

| | gender | age_bucket | urbanicity | prefer_corrugated |
|-----|--------|------------|------------|-------------------|
| 252 | female | 18_to_35 | urban | 1 |
| 253 | female | 18_to_35 | urban | 1 |
| 254 | female | 18_to_35 | urban | 1 |

```
In [15]: #now we need to 'select' for only the 'prefer_corrugated' column
poll[poll["gender"]=="male"]['prefer_corrugated'].head(3)
```

Out[15]:

| | |
|---|---|
| 0 | 1 |
| 1 | 1 |
| 2 | 1 |

Name: prefer_corrugated, dtype: int64

```
In [16]: poll[poll["gender"]=="female"]['prefer_corrugated'].head(3)
```

Out[16]:

| | |
|-----|---|
| 252 | 1 |
| 253 | 1 |
| 254 | 1 |

Name: prefer_corrugated, dtype: int64

```
In [17]: #we'll store these as variables for easier use later
men = poll[poll["gender"]=="male"]['prefer_corrugated']
women = poll[poll["gender"]=="female"]['prefer_corrugated']
```

```
In [18]: #now for use in the proportions test, we need to make a mini dataframe
#counts is the number of successes
#in a binomial setup where it's just 1s and 0s, we can just use sum() to get this
#nobs is the total number of trials, len() can work, though I chose to use counts()
gender_polls = pd.DataFrame({
    "count": [men.sum(), women.sum()], #those that prefer corrugated
    "nobs": [men.count(), women.count()]
}, index=['men', 'women'])
```

```
In [19]: gender_polls
```

Out[19]:

| | count | nobs |
|-------|-------|------|
| men | 1598 | 2462 |
| women | 1345 | 2539 |

```
In [20]: #now we use this to feed into the stats test
#for some odd reason if you say gender_polls.count it will blow up...
#so you have to say gender_polls['count']
statsmodels.stats.proportion.proportions_ztest(gender_polls['count'], gender_polls['nobs'])
#z score, p-value
```

Out[20]:

| | |
|---------------------|-------------------------|
| (8.573032591956961, | 1.0079496366897543e-17) |
|---------------------|-------------------------|

So we can reject the null, if we have a 8.57 z score, and p value with 16 zeros in front.

Let's also see how this would work for a t-test. Two sample mean t test

```
In [22]: from scipy import stats
```

```
In [23]: #you may be surprised, but we can actually use our work from above
#when we defined 'men' and 'women' to be the arrays of 1s and 0s for who preferred corrugated, siloed out for men and women, respectively
stats.ttest_ind(men, women)
```

Out[23]:

| | |
|--|-------------------------------|
| Ttest_indResult(statistic=8.635004923552513, | pvalue=7.790736712355893e-18) |
|--|-------------------------------|

So as we might expect, there's a little difference using the t distribution, but you can still see basically we have a 8.63 t score for whatever the degrees of freedom were and a very low p value.

What this test is saying in both the z and t examples above is that there is a very low chance that we'd see a difference this wide between men and women by a mere random fluctuation.

So we can reject the null, and say there's probably a 'there, there'!

```
In [24]: #btw, just for the heck of it, let me show you what those numbers really were:
gender_polls['proportion'] = round(gender_polls['count'] / gender_polls['nobs'],2)
gender_polls
```

Out[24]:

| | count | nobs | proportion |
|-------|-------|------|------------|
| men | 1598 | 2462 | 0.65 |
| women | 1345 | 2539 | 0.53 |